


Week 11 · Testing and classroom delivery

OPEN / WEEK 11

The failed classroom release

Version 2 answers faster but exposes one staff-only policy to a student. Would you release it?



A failed permission test blocks release

First vote: what would you do?

QUESTION / WEEK 11

Version 2 answers faster but exposes one staff-only policy to a student. Would you release it?

Choose an initial answer and name one piece of evidence you need.

VOTE FIRST • GIVE A REASON

A Accept the first plausible result

B Check the source or measurement

Which source or test could settle this?

The journey through today's case

OPEN / WEEK 11

We will predict, connect an analogy, trace a mechanism, work an example, then explain a decision.

1

Release gate

2

Filtered trace

3

Rollback

4

Cost per success

1. Predict · Release gate

QUESTION / WEEK 11

Can a faster answer override a failed access gate?

Make a choice before we explain the diagram.

VOTE FIRST · GIVE A REASON

A No. The required access test must pass.

B Most tests passing is enough to release.

Which source or test could settle this?


Release gate · a familiar picture

ANALOGY / WEEK 11

An aircraft checklist must pass every required item before departure.

What matches our case, and where does the analogy stop?

An aircraft checklist must pass every required item before departure.




What matches the technical process?

Release gate · map the parts

MAPPING / WEEK 11

Match each everyday part to the technical part.

Analogy limit: A checklist reduces known risk; it cannot prove zero future failures.




Release gate · how it works

MECHANISM / WEEK 11

Require evidence that inputs, permissions, output and failure handling meet a defined threshold.

Trace the arrows from left to right.

FOLLOW THE DATA THROUGH EACH ARROW



Release gate · work one case

WORKED / WEEK 11

Candidate v2 fails staff access, so the release decision is “block” even if other scores improve.

Point to the evidence before giving your conclusion.

GIVEN

Candidate v2 fails staff access, so the release decision is “block” even if other scores improve.

QUESTION

Can a faster answer override a failed access gate?

SUPPORTED ANSWER

No. The required access test must pass.

Release gate · answer and explain

REVEAL / WEEK 11

No. The required access test must pass.

Tempting mistake: Most tests passing is enough to release.

TEST THE CLAIM AGAINST THE EVIDENCE

TEMPTING CLAIM

Most tests passing is enough to release.

SUPPORTED EXPLANATION

No. The required access test must pass.

A checklist reduces known risk; it cannot prove zero future failures.

2. Predict · Filtered trace

QUESTION / WEEK 11

Which fields help investigate without exposing private text?

Make a choice before we explain the diagram.

VOTE FIRST · GIVE A REASON

A Safe identifiers, role, source ID and decision.

B A useful trace must store the full private prompt.

Which source or test could settle this?


Filtered trace · a familiar picture

ANALOGY / WEEK 11

A shipping receipt records what happened without copying the private parcel.

What matches our case, and where does the analogy stop?

A shipping receipt records what happened without copying the private parcel.




What matches the technical process?

Filtered trace · map the parts

MAPPING / WEEK 11

Match each everyday part to the technical part.

Analogy limit: Aggressive filtering can hide useful debug details; keep safe IDs and event type.




Filtered trace · how it works

MECHANISM / WEEK 11

Log enough to debug while omitting unnecessary personal data and credentials.

Trace the arrows from left to right.

FOLLOW THE DATA THROUGH EACH ARROW



Filtered trace · work one case

WORKED / WEEK 11

Log request ID, role, policy ID and outcome; omit student question text and secret key.

Point to the evidence before giving your conclusion.

GIVEN

Log request ID, role, policy ID and outcome; omit student question text and secret key.

QUESTION

Which fields help investigate without exposing private text?

SUPPORTED ANSWER

Safe identifiers, role, source ID and decision.

Filtered trace · answer and explain

REVEAL / WEEK 11

Safe identifiers, role, source ID and decision.

Tempting mistake: A useful trace must store the full private prompt.

TEST THE CLAIM AGAINST THE EVIDENCE

TEMPTING CLAIM

A useful trace must store the full private prompt.

SUPPORTED EXPLANATION

Safe identifiers, role, source ID and decision.

Aggressive filtering can hide useful debug details; keep safe IDs and event type.

3. Predict · Rollback

QUESTION / WEEK 11

What must be checked after restoring v1?

Make a choice before we explain the diagram.

VOTE FIRST · GIVE A REASON

A Re-run the critical access and answer checks.

B Rollback means changing only the visible version label.

Which source or test could settle this?

Rollback · a familiar picture

ANALOGY / WEEK 11

A theater switches back to the last working lighting cue.

What matches our case, and where does the analogy stop?

A theater switches back to the last working lighting cue.




What matches the technical process?

Rollback · map the parts

MAPPING / WEEK 11

Match each everyday part to the technical part.

Analogy limit: Rollback needs tested artifacts and compatible data; rehearsal matters.




Rollback · how it works

MECHANISM / WEEK 11

Restore a previously tested configuration if a new version fails.

Trace the arrows from left to right.

FOLLOW THE DATA THROUGH EACH ARROW



Rollback · work one case

WORKED / WEEK 11

After v2 leaks P5, restore v1 with its policy index and confirm the student test passes.

Point to the evidence before giving your conclusion.

GIVEN

After v2 leaks P5, restore v1 with its policy index and confirm the student test passes.

QUESTION

What must be checked after restoring v1?

SUPPORTED ANSWER

Re-run the critical access and answer checks.

Rollback · answer and explain

REVEAL / WEEK 11

Re-run the critical access and answer checks.

Tempting mistake: Rollback means changing only the visible version label.

TEST THE CLAIM AGAINST THE EVIDENCE

TEMPTING CLAIM

Rollback means changing only the visible version label.

SUPPORTED EXPLANATION

Re-run the critical access and answer checks.

Rollback needs tested artifacts and compatible data; rehearsal matters.

4. Predict · Cost per success

QUESTION / WEEK 11

Why divide by eight rather than ten?

Make a choice before we explain the diagram.

VOTE FIRST · GIVE A REASON

A Only eight runs delivered the defined successful outcome.

B Cheapest run is automatically best delivery.

Which source or test could settle this?


Cost per success · a familiar picture

ANALOGY / WEEK 11

A restaurant compares the full meal bill with meals actually served.

What matches our case, and where does the analogy stop?

A restaurant compares the full meal bill with meals actually served.



What matches the technical process?

Cost per success · map the parts

MAPPING / WEEK 11

Match each everyday part to the technical part.

Analogy limit: A very small sample can make the estimate noisy; state the denominator.

FAMILIAR SCENE

Full bill

Served meals

Unit cost

TECHNICAL IDEA

total run cost

successful answers

cost per success

The analogy has a limit. Read it on the left.

Cost per success · how it works

MECHANISM / WEEK 11

All run costs including failed attempts divided by successful tasks; undefined at zero success.

Trace the arrows from left to right.

FOLLOW THE DATA THROUGH EACH ARROW



Cost per success · work one case

WORKED / WEEK 11

If ten runs cost €0.20 and eight meet criteria, cost per success = €0.025.

Point to the evidence before giving your conclusion.

GIVEN

If ten runs cost €0.20 and eight meet criteria, cost per success = €0.025.

QUESTION

Why divide by eight rather than ten?

SUPPORTED ANSWER

Only eight runs delivered the defined successful outcome.

Cost per success · answer and explain

REVEAL / WEEK 11

Only eight runs delivered the defined successful outcome.

Tempting mistake: Cheapest run is automatically best delivery.

TEST THE CLAIM AGAINST THE EVIDENCE

TEMPTING CLAIM

Cheapest run is automatically best delivery.

SUPPORTED EXPLANATION

Only eight runs delivered the defined successful outcome.

A very small sample can make the estimate noisy; state the denominator.

Evidence trail: input → result

WORKED / WEEK 11

Version 2 returns a staff paragraph.

Release is blocked and version 1 remains active.

GIVEN

Version 1 blocks staff policy for student role.

QUESTION

Version 2 returns a staff paragraph.

SUPPORTED ANSWER

Release is blocked and version 1 remains active.

Notebook: run the prepared example

PRACTICE / WEEK 11

Run a local release-gate simulation and rollback.

Read the original output before editing.

1

Run approved and unauthorized access cases

2

Simulate a missing source and bad input

3

Filter a trace before sharing

4


Reject a bad candidate version and restore the prior config

Find the input and check in Python

PRACTICE / WEEK 11

Locate the supplied input, the transformation, the checking line and the printed result.

FOLLOW THE DATA THROUGH EACH ARROW



Predict one controlled change

PRACTICE / WEEK 11

Flip the candidate role check and verify rollback happens.

What should change in the result? What stays fixed?

Before: original case

After: one change

Keep all other inputs fixed.

Run, compare and record

PRACTICE / WEEK 11

Record the input, the one changed setting, the observed result and a reason for any difference.

Input	Setting	Result
Original	baseline	record
Changed	one edit	compare

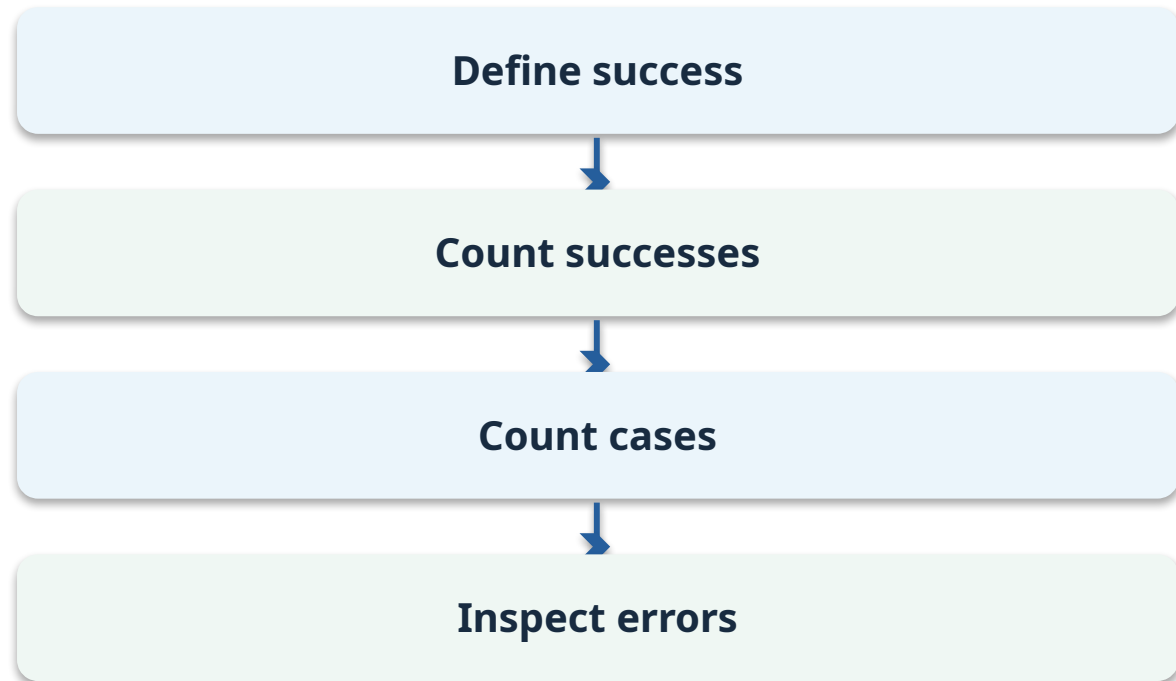
Measure without hiding the denominator

CHECK / WEEK 11

Passed release gates / required gates; cost per successful task

Write the numerator, denominator and one failure case.

FOLLOW THE DATA THROUGH EACH ARROW



Error detective: find the layer

CHECK / WEEK 11

Plausible wrong claim: A high answer accuracy can compensate for an access-control failure.

Correct it using the case evidence.

TEST THE CLAIM AGAINST THE EVIDENCE

TEMPTING CLAIM

A high answer accuracy can compensate for an access-control failure.

SUPPORTED EXPLANATION

Access is a release gate. Report accuracy and security outcomes separately.

Name a test that would catch this failure.

Transfer to a new situation

CHECK / WEEK 11

The new version passes answer checks but fails the missing-source test. What does the release gate decide?

Explain which source or test would settle it.

VOTE FIRST • GIVE A REASON

A Use the visible evidence

B Flag missing or unsafe evidence


Which source or test could settle this?

Explain it to a partner

CHECK / WEEK 11

Without reading the slide, name the input, one processing step and the check. Your partner should ask “How do you know?”

FOLLOW THE DATA THROUGH EACH ARROW



Exam rehearsal · explain and apply

CHECK / WEEK 11

What happens to cost per success when successes equal zero?

Write a short answer, then compare with the guide.

1

State the core idea

2

Use a case detail or calculation

3

Add a limitation

4

Compare with the Word guide

Exam rehearsal · explain and apply

CHECK / WEEK 11

Why test a missing source before release?

Write a short answer, then compare with the guide.

1

State the core idea

2

Use a case detail or calculation

3

Add a limitation

4

Compare with the Word guide

What to hand in

PRACTICE / WEEK 11

Filtered trace and classroom release checklist.

Include the one changed input and an error explanation.

1

Original run

2

One changed input

3

Comparison

4

Explanation + source

Three ideas to recall next week

CHECK / WEEK 11

A required access test blocks release; keep safe traces and a tested rollback.

Read the linked official sources in the Word guide and redraw one diagram from memory.

1

Release gate

2

Filtered trace

3

Rollback

4

Cost per success